



# Clustering

CE-477: Machine Learning - CS-828: Theory of Machine Learning  
Sharif University of Technology  
Fall 2024


Fatemeh Seyyedsalehi

# Two views on clustering

- ▶ We have a set of unlabeled data points  $\{\mathbf{x}^{(i)}\}_{i=1}^N$
- ▶ Clusters:
  - ▶ Similarity or distance based definition
    - ▶ high intra-cluster similarity
    - ▶ low inter-cluster similarity
  - ▶ Density-based definition:
    - ▶ Clusters are regions of high density that are separated from one another by regions of low density

# Clustering Purpose

- ▶ Automatically organizing data. Ex. Compression, quantization, visualization

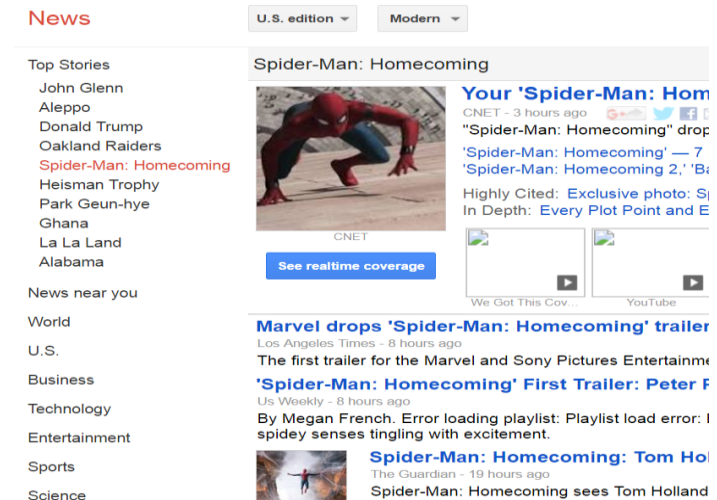


- ▶ Knowledge discovery from data: As a tool to **understand the hidden structure** in data or to **group** them
  - ▶ To gain insight into the structure of the data (prior to classifier design)
  - ▶ Provides information about the internal structure of the data

# Clustering Applications


- ▶ Cluster news articles or web pages or search results by topic.

- ▶ Google news




- ▶ Cluster users of social networks by interest (community detection).

Facebook  
network



Twitter  
network





# Clustering methods we will discuss

- ▶ Objective based clustering

- ▶ Construct various partitions and then evaluate them by some criterion
  - ▶ K-means
  - ▶ EM-style algorithm for clustering for mixture of Gaussians (in the next lecture)

- ▶ Hierarchical clustering

- ▶ Create a hierarchical decomposition of the set of objects using some criterion

# Partitioning Algorithms: Basic Concept

---

- ▶ Construct a partition of a set of  $N$  objects into a set of  $K$  clusters
  - ▶ The number of clusters  $K$  is given in advance
  - ▶ Each object belongs to **exactly one** cluster in hard clustering methods
- ▶ K-means is the most popular partitioning algorithm
  - ▶ K-means was proposed near 60 years ago
  - ▶ Thousands of clustering algorithms have been published since then
  - ▶ However, K-means is still widely used.

# K-means Clustering

- ▶ **Input:** a set  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}$  of data points (in a  $d$ -dim feature space) and an integer  $K$
- ▶ **Output:** a set of  $K$  representatives  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K \in \mathbb{R}^d$  as the cluster representatives
  - ▶ data points are assigned to the clusters according to their distances to  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$ 
    - ▶ Each data is assigned to the cluster whose representative is nearest to it
- ▶ **Objective:** choose  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$  to minimize:

$$\sum_{i=1}^N \min_{j \in 1, \dots, K} d(\mathbf{x}^{(i)}, \mathbf{c}_j)$$



# Euclidean k-means Clustering

- ▶ **Input:** a set  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}$  of data points (in a  $d$ -dim feature space) and an integer  $K$
- ▶ **Output:** a set of  $K$  representatives  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K \in \mathbb{R}^d$  as the cluster representatives
  - ▶ data points are assigned to the clusters according to their distances to  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$ 
    - ▶ Each data is assigned to the cluster whose representative is nearest to it
- ▶ **Objective:** choose  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$  to minimize:

$$\sum_{i=1}^N \min_{j \in 1, \dots, K} \|\mathbf{x}^{(i)} - \mathbf{c}_j\|^2$$

each point assigned to its closest cluster representative

# Euclidean k-means Clustering: Computational Complexity

- ▶ NP hard: even for  $k = 2$  or  $d = 2$
- ▶ For  $k=1$ :  $\min_{\mathbf{c}} \sum_{i=1}^N \|\mathbf{x}^{(i)} - \mathbf{c}\|^2$ 
  - ▶  $\mathbf{c} = \boldsymbol{\mu} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$
- ▶ For  $d = 1$ , dynamic programming in time  $O(N^2 K)$ .

## Common Heuristic in Practice: The Lloyd's method

► Input: A set  $\mathcal{X}$  of  $N$  data points  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}$  in  $\mathbb{R}^d$

► **Initialize** centers  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K \in \mathbb{R}^d$  in any way.

► **Repeat** until there is no further change in the cost.

► For each  $j$ :  $\mathcal{C}_j \leftarrow \{\mathbf{x} \in \mathcal{X} \mid \text{where } \mathbf{c}_j \text{ is the closest center to } \mathbf{x}\}$

► For each  $j$ :  $\mathbf{c}_j \leftarrow \text{mean of members of } \mathcal{C}_j$

Holding centers  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$  fixed

Find optimal assignments  $\mathcal{C}_1, \dots, \mathcal{C}_K$  of data points to clusters











Holding cluster assignments  $\mathcal{C}_1, \dots, \mathcal{C}_K$  fixed

Find optimal centers  $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K$

$$\mathbf{c}_j = \frac{1}{|\mathcal{C}_j|} \sum_{\mathbf{x}^{(i)} \in \mathcal{C}_j} \mathbf{x}^{(i)}$$

## Assigning data to clusters

## Updating means



[Bishop]

## Intra-cluster similarity view

- ▶ k-means optimizes intra-cluster similarity:

$$J(\mathcal{C}) = \sum_{j=1}^K \sum_{\mathbf{x}^{(i)} \in \mathcal{C}_j} \|\mathbf{x}^{(i)} - \mathbf{c}_j\|^2$$


$$\mathbf{c}_j = \frac{1}{|\mathcal{C}_j|} \sum_{\mathbf{x}^{(i)} \in \mathcal{C}_j} \mathbf{x}^{(i)}$$

$$\sum_{\mathbf{x}^{(i)} \in \mathcal{C}_j} \|\mathbf{x}^{(i)} - \mathbf{c}_j\|^2 = \frac{1}{2|\mathcal{C}_j|} \sum_{\mathbf{x}^{(i)} \in \mathcal{C}_j} \sum_{\mathbf{x}^{(i')} \in \mathcal{C}_j} \|\mathbf{x}^{(i)} - \mathbf{x}^{(i')}\|^2$$

the average distance to members of the same cluster

# K-means: Convergence


- ▶ It always converges.
  - ▶ The cost always drop
  - ▶ There is only a finite number of partitioning
- ▶  $K$ -means algorithm reaches a state in which clustering doesn't change.
  - ▶ Reassignment stage monotonically decreases  $J$  since each vector is assigned to the closest centroid.
  - ▶ Centroid update stage also for each cluster minimizes the sum of squared distances of the assigned points to the cluster from its center.




# Local optimum problem

- ▶ It always converges
  - ▶ But it may converge at a local optimum
- ▶ Colored points show initial centroids, two resulted partitioning for two different initial centroids

Local optimum






Best partitioning or  
global optimum



# Local optimum problem


- ▶ It always converges
  - ▶ But it may converge at a local optimum
- ▶ This bad performance, can happen even with well separated Gaussian clusters.





# Local optimum problem

- ▶ It always converges
  - ▶ But it may converge at a local optimum
- ▶ This bad performance, can happen even with well separated Gaussian clusters.





# The Lloyd's method: Initialization

- ▶ Initialization is crucial (how fast it converges, quality of clustering)
  - ▶ Random centers from the data points
    - ▶ Multiple runs and select the best ones
  - ▶ Initialize with the results of another method
  - ▶ Select good initial centers using a heuristic
    - ▶ Furthest traversal
    - ▶ K-means ++ (works well and has provable guarantees)

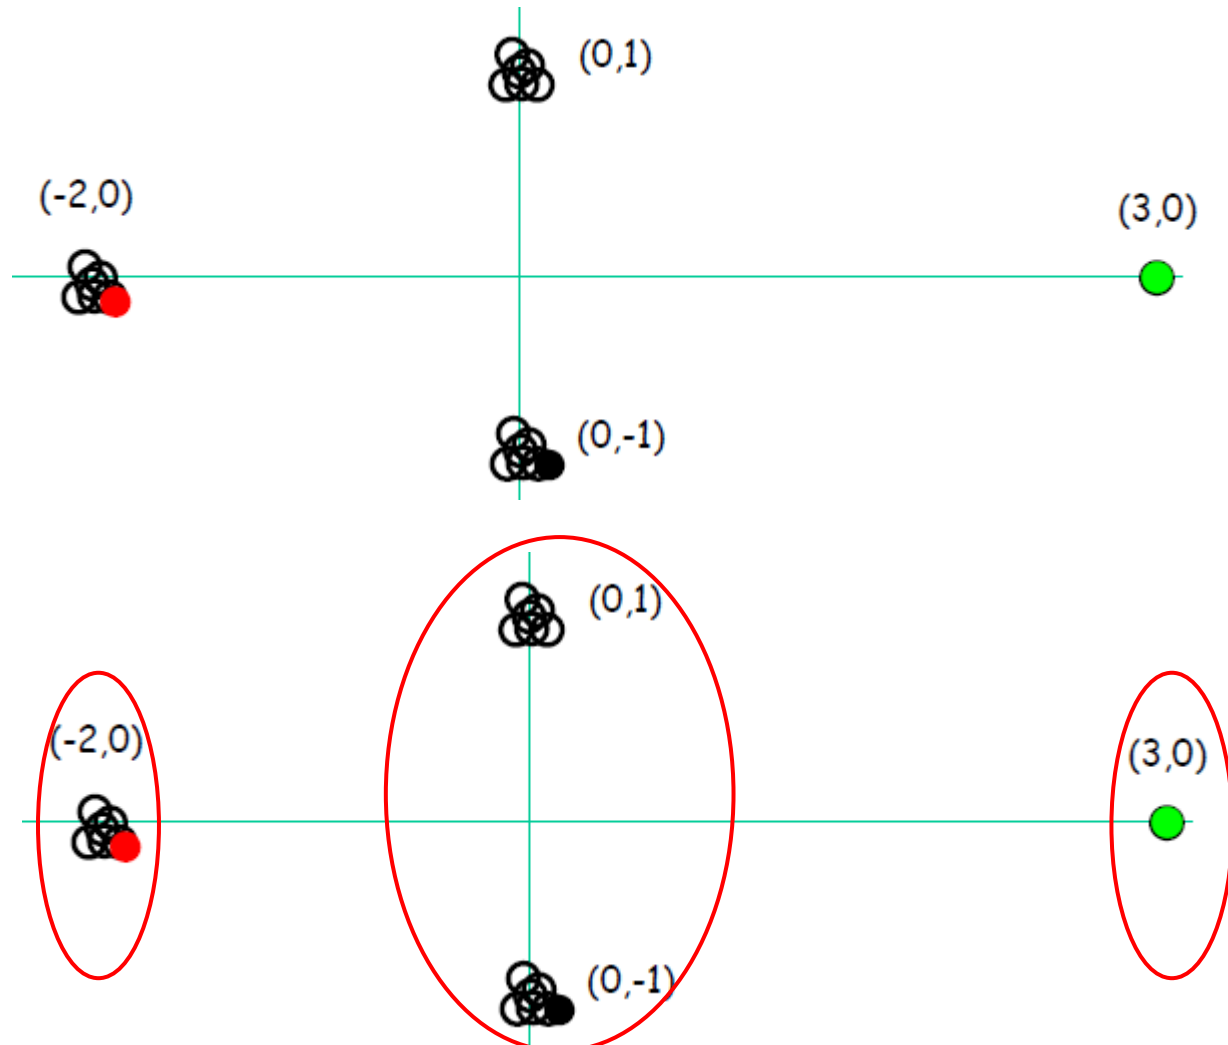
# Initialization Idea: Furthest Point Heuristic

- ▶ Choose  $\mathbf{c}_1$  arbitrarily (or at random).
- ▶ For  $j = 2, \dots, K$ 
  - ▶ Select  $\mathbf{c}_j$  among datapoints  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}$  that is farthest from previously chosen  $\mathbf{c}_1, \dots, \mathbf{c}_{j-1}$
- ▶ Good for our previous example



# Initialization Idea: Furthest Point Heuristic

- ▶ However, it is sensitive to outliers



# K-means++ Initialization: D2 sampling

[D. Arthur and S. Vassilvitskii, 2007]

- ▶ Combine random initialization and furthest point initialization ideas
- ▶ Let the probability of selection of the point be proportional to the distance between this point and its nearest center.
  - ▶ probability of selecting of  $\mathbf{x}$  is proportional to  $D^2(\mathbf{x}) = \min_{k < j} \|\mathbf{x} - \mathbf{c}_k\|^2$ .

- ▶ Choose  $\mathbf{c}_1$  arbitrarily (or at random).
- ▶ For  $j = 2, \dots, K$ 
  - ▶ Select  $\mathbf{c}_j$  among data points  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}$  according to the distribution:

$$\Pr(\mathbf{c}_j = \mathbf{x}^{(i)}) \propto \min_{k < j} \|\mathbf{x}^{(i)} - \mathbf{c}_k\|^2$$

# How Many Clusters?

- ▶ Number of clusters  $k$  is given in advance in the k-means algorithm
  - ▶ However, finding the “right” number of clusters is a part of the problem
- ▶ Tradeoff between having better focus within each cluster and having too many clusters

# How Many Clusters?

- ▶ Heuristic:

- ▶ Find large gap between  $k - 1$ -means cost and  $k$ -means cost.
- ▶ “knee finding” or “elbow finding”.


- ▶ Hold-out validation/cross-validation on auxiliary task (e.g., supervised learning task).

- ▶ Optimization problem: penalize having lots of clusters

- ▶ some criteria can be used to automatically estimate  $k$ 
  - ▶ Penalize the number of bits you need to describe the extra parameter

$$J'(\mathcal{C}) = J(\mathcal{C}) + |\mathcal{C}| \times \log N$$

- ▶ Hierarchical clustering





# K-means summary

- ▶ Relatively efficient:  $O(tKNd)$ , where  $t$  is the number of iterations.
  - ▶  $K$ -means typically converges quickly
    - ▶ Usually  $t \ll n$ .
  - ▶ Exponential # of rounds in the worst case [Andrea Vattani 2009].
- ▶ Limitations
  - ▶ Need to specify  $K$ , the *number* of clusters, in advance
  - ▶ Often terminates at a *local optimum*.
    - ▶ Initialization is important.
  - ▶ Not suitable to discover clusters with arbitrary shapes
  - ▶ Works for numerical data. What about categorical data?
  - ▶ Noise and outliers can be considerable trouble to  $K$ -means





# k-means summary

- ▶ In general, k-means is unable to find clusters of arbitrary shapes, sizes, and densities
  - ▶ Except to very distant clusters



# Hierarchical Clustering

- ▶ Hierarchical Clustering: Clusters contain sub-clusters and sub-clusters themselves can have sub-sub-clusters, and so on
  - ▶ Several levels of details in clustering
- ▶ A hierarchy might be more natural.
  - ▶ Different levels of granularity



# Hierarchical Clustering

- ▶ Agglomerative (bottom up):

- ▶ Starts with each data in a separate cluster
- ▶ Repeatedly joins the closest pair of clusters, until there is only one cluster (or other stopping criteria).
  - ▶ Different definition for distance between clusters



- ▶ Divisive (top down):

- ▶ Starts with the whole data as a cluster
- ▶ Repeatedly divide data in one of the clusters until there is only one data in each cluster (or other stopping criteria).

# Agglomerative (bottom up):

1. Initially, each instance forms a cluster
2. While there are more than one cluster  
Pick the two closest one  
Merge them into a new cluster


Height represents the distance at which the merge occurs



## Agglomerative (bottom up):


- ▶ Having a distance measure  $dist_{SL}(x, y)$  on a pair of objects, many variants exist to define distances between pair of clusters
  - ▶ **Single-link**
    - ▶ Minimum distance between different pairs of data
  - ▶ **Complete-link**
    - ▶ Maximum distance between different pairs of data
  - ▶ **Centroid (Ward's)**
    - ▶ Distance between centroids (centers of gravity)
  - ▶ **Average-link**
    - ▶ Average distance between pairs of elements

# Distances between Cluster Pairs




**Single-link**

$$dist_{SL}(\mathcal{C}_i, \mathcal{C}_j) = \min_{x \in \mathcal{C}_i, x' \in \mathcal{C}_j} dist(x, x')$$




**Complete-link**

$$dist_{CL}(\mathcal{C}_i, \mathcal{C}_j) = \max_{x \in \mathcal{C}_i, x' \in \mathcal{C}_j} dist(x, x')$$



**Ward's**


$$dist_{Ward}(\mathcal{C}_i, \mathcal{C}_j) = \frac{|\mathcal{C}_i| |\mathcal{C}_j|}{|\mathcal{C}_i| + |\mathcal{C}_j|} dist(\mathbf{c}_i, \mathbf{c}_j)$$



**Average-link**


$$dist_{AL}(\mathcal{C}_i, \mathcal{C}_j) = \frac{1}{|\mathcal{C}_i \cup \mathcal{C}_j|} \sum_{x \in \mathcal{C}_i \cup \mathcal{C}_j} \sum_{x' \in \mathcal{C}_i \cup \mathcal{C}_j} dist(x, x')$$

# Single-Link




Single linkage can produce long stretched clusters.

Dendrogram




# Complete Link



One way to think of it: keep max diameter as small as possible at any level.

Complete linkage prefers compact clusters.






## Ward's method

- ▶ The distances between centers of the two clusters (weighted to consider sizes of clusters too):

$$\text{dist}_{\text{Ward}}(\mathcal{C}_i, \mathcal{C}_j) = \frac{|\mathcal{C}_i| |\mathcal{C}_j|}{|\mathcal{C}_i| + |\mathcal{C}_j|} \text{dist}(\mathbf{c}_i, \mathbf{c}_j)$$

- ▶ Merge the two clusters such that the increase in k-means cost is as small as possible.
- ▶ Works well in practice.




# Computational Complexity

- ▶ In the first iteration, all HAC methods compute similarity of all pairs of  $N$  individual instances which is  $O(N^2)$  similarity computation.
- ▶ In each  $N - 1$  merging iterations, compute the distance between the most recently created cluster and all other existing clusters.
- ▶ If done naively  $O(N^3)$  but if done more cleverly  $O(N^2 \log N)$

# Dendrogram: Hierarchical Clustering

- ▶ Clustering obtained by cutting the dendrogram at a desired level
  - ▶ Cut at a pre-specified level of similarity
  - ▶ where the gap between two successive combination similarities is largest
  - ▶ select the cutting point that produces  $K$  clusters



# K-means vs. Hierarchical

- ▶ Time cost:
  - ▶ K-means is usually fast while hierarchical methods do not scale well
- ▶ Human intuition
  - ▶ Hierarchical structure provides more natural output compatible with human intuition in some domains
- ▶ Local minimum problem
  - ▶ It is very common for k-means
  - ▶ Hierarchical methods like any heuristic search algorithms also suffer from local optima problem.
    - ▶ Since they can never undo what was done previously and greedily merge clusters
- ▶ Choosing of the number of clusters
  - ▶ There is no need to specify the number of clusters in advance for hierarchical methods